

SIMON OKOSODO

Junior Data Scientist | Machine Learning & Applied Analytics

Lagos, Nigeria | simonokosodopatrick@gmail.com | linkedin.com/in/simon-okosodo-ai | github.com/simon-okosodo-ds | simon-okosodo-ds.netlify.app

PROFESSIONAL SUMMARY

Data science intern with hands-on experience turning a real, messy search-performance dataset into a validated, honestly-evaluated machine learning solution. Built a content-prioritization and diagnosis system that nearly tripled review precision over an existing rule-based baseline (Precision@50: 0.24 to 0.74), using client-grouped validation and a formal leakage-hunting process rather than an inflated, unchecked score. Comfortable working the full pipeline end to end: framing a business decision, writing and verifying a data contract, engineering leakage-safe features, training and validating models honestly, and communicating results in language a non-technical stakeholder can trust. Actively building AI-agent and generative-AI fluency alongside applied ML (see Certifications). Background in Chemical Engineering underpins a systems-level, assumption-checking approach to data work.

CORE SKILLS

- **Languages & Tools:** Python (pandas, scikit-learn, NumPy), SQL, DuckDB, Git/GitHub, Jupyter/Colab
- **Machine Learning:** Classification & ranking models, logistic regression, random forests, feature engineering, precision@K evaluation, baseline-vs-model comparison
- **Honest Validation:** Client-grouped & time-aware train/test splits, leakage detection and removal, base-rate-aware reporting, claim-ladder discipline
- **Data & Analytics Sources:** Google Search Console, Google Analytics (GA4), BigQuery-derived warehouse data, DuckDB over remote Parquet at multi-million-row scale
- **AI Agents & Generative AI:** Agentic engineering, prompt engineering, generative AI foundations, AI-assisted / vibe-coding workflows
- **AI-Assisted Workflow Discipline:** Cross-model evaluation (Claude vs. ChatGPT) on real debugging tasks, disciplined human-in-the-loop review of all AI-drafted code and claims

CERTIFICATIONS

- Google Antigravity: Agentic Application Development & AI Engineering — Google Developer Program (Jun 2026)
- Foundations in Generative AI — IBM SkillsBuild (Jun 2026)
- Prompt Engineering: Shaping Better AI Responses — IBM SkillsBuild (Jul 2026)
- 5-Day AI Agents: Intensive Vibe Coding Course — Kaggle / Google (Jul 2026)

EXPERIENCE

Data Science Intern — FlyRank

(Jul 2026 – Present)

Remote — content intelligence / applied ML

- Framed and led a self-directed capstone lane ("Diagnosis-First Content Triage") on a real, pseudonymized search-performance dataset (30,000+ pages, 32+ clients, scaling to a ~79M-row warehouse), identifying a previously invisible segment (3.7% of pages) that the company's existing rule-based system missed entirely.
- Reproduced and validated the project's core business case: a trained classifier lifts review precision from 0.24 to 0.74 (Precision@50) versus the existing hand-written rule — nearly tripling the share of editor reviews that are actually worthwhile at the same review capacity.
- Engineered leakage-safe diagnostic features built entirely from raw, observed signals — explicitly excluding the system's own internal decision flags to avoid circular, self-confirming results.
- Applied a formal leakage-hunting process and client-grouped validation split, deliberately demonstrating a label-derived leakage trap (a near-perfect but meaningless score) and correcting it before trusting any result.

- Wrote and iteratively engineered prompts across a documented six-version ladder for a real production debugging task, then cross-validated outputs between Claude and ChatGPT.
- Communicated findings using a strict evidence-to-claim discipline (observed / directional / decision-support language only), avoiding unsupported or inflated claims in all written analysis.

Independent Machine Learning Project — Customer Churn Prediction — Self-Directed

(Aug 2026)

- Built and honestly validated a churn model on a real public dataset (IBM Telco Customer Churn, 7,043 customers, 26.6% base churn rate) — same end-to-end method used at FlyRank: transparent baseline rule first, then model comparison, same split and metric.
- Baseline rule (month-to-month contract + tenure under 12 months) scored 0.491 precision; Logistic Regression reached 0.643 precision and 0.835 AUC, meaningfully beating both the baseline and a Random Forest (0.625 precision, 0.815 AUC) — reported the simpler model's win honestly rather than defaulting to the more complex one.
- Identified TotalCharges, MonthlyCharges, and tenure as the top predictive features.

EDUCATION & CONTINUING STUDY

B.Eng., Chemical Engineering — University of Benin (UNIBEN), Nigeria

- Coursework focus: process optimization, numerical methods, systems engineering.

Data Science & Machine Learning Professional Program — GDSC ML Institute, Ghana

- Focus: gradient boosting methods, feature engineering, and model deployment fundamentals.

AI Fluency: Framework & Foundations — Anthropic Academy

- Structured coursework in applied AI collaboration, prompt engineering, and responsible AI-assisted workflows, completed alongside the FlyRank internship.

WHY THIS PROFILE

- Proven ability to work a real, messy, imperfect dataset end to end — framing the business decision, quantifying the cost of getting it wrong, and validating honestly enough to trust the result.
- Demonstrated discipline around leakage and overclaiming — treats an unbelievably good score as a bug to investigate, not a result to celebrate, and reports when a simpler model wins.
- Actively building agentic-AI and generative-AI fluency (4 verified certifications, 2026), on top of core ML skills — fluent in using AI as a force multiplier while keeping human judgment in the loop.